

TAKE HOME TEST 3 OPTIMIZATION

CSC 343 FALL 2017

DECEMBER 4, 2017

Jeter Gutierrez

TABLE OF CONTENTS**1 OBJECTIVE pg. 2****2. DOT PRODUCT ON LINUX 64 BITS pg. 2-13****2.1 DOT PRODUCT USING INDEXES ON LINUX 64 BIT. pg. 3-5****2.2 DOT PRODUCT USING OPTIMIZED INDEX ON LINUX 64 BIT pg. 5-8****2.3 DOT PRODUCT USING POINTERS ON LINUX 64 BIT pg. 8-10****2.4 DOT PRODUCT USING OPTIMIZED POINTER ON LINUX 64 BIT pg. 10-13****3. DOT PRODUCT ON LINUX 32 BIT pg. 13-25****3.1 DOT PRODUCT USING INDEXES ON LINUX 32 BIT. pg. 13-16****3.2 DOT PRODUCT USING OPTIMIZED INDEX ON LINUX 32 BIT pg. 16-19****3.3 DOT PRODUCT USING POINTERS ON LINUX 32 BIT pg. 19-21****3.4 DOT PRODUCT USING OPTIMIZED POINTER ON LINUX 32 BIT pg. 21-25****4. DOT PRODUCT ON WINDOWS 32 BIT pg. 25-38****4.1 DOT PRODUCT USING INDEXES ON WINDOWS 32 BIT. pg. 25-29****4.2 DOT PRODUCT USING OPTIMIZED INDEX ON WINDOWS 32 BIT pg. 29-33****4.3 DOT PRODUCT USING POINTERS ON WINDOWS 32 BIT pg. 33-35****4.4 DOT PRODUCT USING OPTIMIZED POINTER ON WINDOWS 32 BIT pg. 35-38****5. GRAPHS pg. 38-42****5.1 INDEX FOR DIFFERENT VALUES NORMAL AND OPTIMIZED ON ALL 3
PLATFORMS pg. 38-40****5.2 POINTER FOR DIFFERENT VALUES NORMAL AND OPTIMIZED ON ALL 3
PLATFORMS pg. 40-42****6. CONCLUSION pg. 42**

1 OBJECTIVE: The purpose of this assignment is to test how a Windows 32-bit compiler, Linux 64-bit compiler, and Linux 32-bit compiler handle performing dot product using pointers and indexing.

We will be using Visual studio to run these examinations on Windows 32-bit compiler, we will be using GDB and GCC to test the both Linux versions. In both the case of running the examination using array indexing or pointers we will be performing Dot product. When we call the Dot product function the function will be written in assembly while the function call will be made using a CPP file. The procedure for every case will be to write the code for performing Dot product in C++ on each platform, after that we will then generate the assembly code for the function. Once we have generated the assembly code for the function we will then call the from the main file and measure its performance, we will do this for different lengths of arrays. After we have measured the performance for both indexing and pointer Dot product computations we will then optimize the code on our own and measure performance again, doing this for all 3 compiling platforms.

In this report we will be performing the following examinations:

- 1) DOT PRODUCT USING INDEX ON LINUX 64 BITS
- 2) DOT PRODUCT USING POINTERS ON LINUX 64 BIT
- 3) DOT PRODUCT USING INDEX ON LINUX 32 BIT
- 4) DOT PRODUCT USING POINTER ON LINUX 32 BIT
- 5) DOT PRODUCT USING INDEX ON WINDOWS 32 BIT
- 6) DOT PRODUCT USING POINTER ON WINDOWS 32 BIT

2. DOT PRODUCT ON LINUX 64 BITS

In this section we will be performing the dot product using index arithmetic for the arrays. The process is as follows, we will write the code that will perform the dot product using indexes in C++, then we will generate the compiler assembly code for the instruction, after we have done that we will optimize the compiles assembly code. With the assembly code not optimized and with it optimized we as well

will be timing it and comparing the values that we get. With the values that we get for optimized times and not optimized times we will see the difference. We will also be testing this for different sized arrays.

```

2  #include <stdio.h>
3  #include <time.h>
4  #include <stdlib.h>
5  #include <stdint.h>
6  // #include <iostream>
7  using namespace std;
8
9  // void DotProductUsingIndex(int Array1[], int Array2[], int size);
10 void DotProductUsingPointer(int *Array1, int *Array2, int size);
11 #define SIZE 10000000
12 static int FIRSTARRAY[SIZE];
13 static int SECONDARRAY[SIZE];
14 main(int argc, char **argv){
15     uint64_t sum = 0;
16     for(int i=0; i<SIZE; i++){
17         FIRSTARRAY[i]=7;
18         SECONDARRAY[i]=5;
19     }
20     for(int i = 0; i < 10; i++){
21         timespec start, end;
22         clock_gettime(CLOCK_MONOTONIC, &start);
23         // DotProductUsingIndex(Array1, Array2, SIZE);
24         DotProductUsingPointer(Array1, Array2, SIZE);
25         clock_gettime(CLOCK_MONOTONIC, &end);
26         sum += (end.tv_nsec - start.tv_nsec);
27     }
28     printf("size = %i, average is %f ms \n", SIZE, (double)(sum/10.0));
29     return 0;
30 }

```

Figure 1: C++ code that will be used to measure the time that it takes to calculate the dot product using index normally, pointers normally, index optimized, and pointers optimized.

2.1 DOT PRODUCT USING INDEXES ON LINUX 64 BIT.

In this section I will be computing the dot product using GCC generated assembly code on the Linux 64-bit compiler on Ubuntu. To generate the code what I need to do is use the flag -S with the GCC compiler. After it is loaded. I will link the assembly code to the main file that will be used to time how long it takes to perform the dot product using indexes on arrays of sizes; 10, 100, 1000, 10000, 100000, 1000000, and 10000000. Then I will measure and record the time it takes and store the data to later use it on Graph to compare to its optimized code.

```

1 void DotProductUsingIndex(int A[],int B[],int size ){
2     int result = 0;
3     for(int i = 0; i < size; i++)
4         result += ( A[i] * B[i]);
5 }

```

Figure 2: C++ code that will be used to calculate the Dot Product on the Linux 64-bit compiler using indexes.

```

1 .file "DOTPRODUCTINDEX.cpp"
2 .text
3 .globl _Z20DotProductUsingIndexPiS_i
4 .type _Z20DotProductUsingIndexPiS_i, @function
5 _Z20DotProductUsingIndexPiS_i:
6 .LFB0:
7 .cfi_startproc
8 pushq %rbp
9 .cfi_def_cfa_offset 16
10 .cfi_offset 6, -16
11 movq %rsp, %rbp
12 .cfi_def_cfa_register 6
13 movq %rdi, -24(%rbp)
14 movq %rsi, -32(%rbp)
15 movl %edx, -36(%rbp)
16 movl $0, -8(%rbp)
17 movl $0, -4(%rbp)
18 .L3:
19 movl -4(%rbp), %eax
20 cmpl -36(%rbp), %eax
21 jge .L4
22 movl -4(%rbp), %eax
23 cltq
24 leaq 0(%rax,4), %rdx
25 movq -24(%rbp), %rax
26 addq %rdx, %rax
27 movl (%rax), %edx

```

Figure 3: Compiler generated assembly code for running dot product using indexes on Linux 64-bit compiler. We can see that the file is being compiled into assembly code from the file Dotproduct.cpp which has the instructions written in C++ that were generated into assembly.

These are the instructions that use the registers to look at the values that will be added to multiply and add into the compiler. Some of the instructions are redundant because they are repetitive and some of them are unnecessary because they make the process take even longer than they should take if the code were written more efficiently.

```

28     movl    -4(%rbp), %eax
29     cltq
30     leaq    0(%rax,4), %rcx
31     movq    -32(%rbp), %rax
32     addq    %rcx, %rax
33     movl    (%rax), %eax
34     imull   %edx, %eax
35     addl    %eax, -8(%rbp)
36     addl    $1, -4(%rbp)
37     jmp     .L3
38 .L4:
39     nop
40     popq    %rbp
41     .cfi_def_cfa 7, 8
42     ret
43     .cfi_endproc
44 .LFE0:
45     .size   _Z20DotProductUsingIndexPiS_i, .-_Z20DotProductUsingIndexPiS_i
46     .ident   "GCC: (Ubuntu 5.4.0-6ubuntu1-14.04.4) 5.4.0 20160609"
47     .section .note.GNU-stack,"",@progbits

```

Figure 4: Compiler generated assembly code that is also used for calculating dot product using indexes on Linux 64-bit compiler. The final instructions of deallocating the memory from the stack pointer and from the stack completely.

LINUX 64 BIT TIME VALUES FOR INDEX							
N=	10	100	1000	10000	100000	1000000	10000000
TRIAL	0.244	1.005	8.369	82.227	825.2497	6254.3498	61412.5957
TRIAL	0.256	0.989	8.326	81.788	819.9741	6268.9691	61373.9987
TRIAL	0.251	0.986	8.357	81.891	827.4812	6251.973	61272.2451
TRIAL	0.247	0.986	8.357	81.772	711.6758	6221.9977	61109.8625
TRIAL	0.243	0.993	8.363	81.714	814.9616	6226.9499	61408.5084
TRIAL	0.247	1.024	8.328	81.695	801.4223	6253.0208	61264.6651
TRIAL	0.244	0.985	8.294	81.71	724.8512	6252.2141	61208.5171
TRIAL	0.247	0.995	8.348	81.816	812.2048	6267.7566	61215.4541
TRIAL	0.249	0.984	8.323	81.742	834.4072	6215.4807	61470.7489
TRIAL	0.252	1.012	8.326	81.891	749.5387	6264.4722	61201.0715

Figure 5: Table that was recorded of the data when running the dot product using the assembly code generated by Linux 64-bit compiler for the different values of N for the different sizes of the arrays. This information will later be graphed and compared to other sections of this project.

2.2 DOT PRODUCT USING OPTIMIZED INDEX ON LINUX 64 BIT

In this section I will be computing the dot product using GCC generated assembly code that I have optimized on the Linux 32-bit compiler on a Raspberry Pi To generate the code what I need to do is use the flag -S with the GCC compiler. After it is loaded. I will optimize it and link the assembly code to

the main file that will be used to time how long it takes to perform the dot product using indexes optimized on arrays of sizes; 10, 100, 1000, 10000, 100000, 1000000, and 10000000. Then I will measure and record the time it takes and store the data to later use it on Graph to compare to its normal assembly code.

```

1 void DotProductUsingIndex(int A[],int B[],int size ){
2     int result = 0;
3     for(int i = 0; i < size; i++)
4         result += ( A[i] * B[i]);
5 }

```

Figure 6: C++ code that will be used to generate the assembly code using the Linux 64-bit compiler on Ubuntu and will then be optimized by me. After I optimize the code I will then link it to the timing file that will be used, the main file that will be used to measure the time that it takes to run the Dot Product using my optimized assembly code. I will be doing this for different values of sized arrays.

```

1 .file "index.cpp"
2 .text
3 .globl _Z20DotProductUsingIndexPiS_i
4 .type _Z20DotProductUsingIndexPiS_i, @function
5 _Z20DotProductUsingIndexPiS_i:
6 .LFB0:
7 .cfi_startproc
8 pushq %rbp
9 .cfi_def_cfa_offset 16
10 .cfi_offset 6, -16
11 movq %rsp, %rbp
12 .cfi_def_cfa_register 6
13 movq %rdi, -24(%rbp)
14 movq %rsi, -32(%rbp)
15 movl %edx, -36(%rbp)
16 movl $0, -8(%rbp)
17 movl $0, -4(%rbp)
18 .L3:
19 movl -4(%rbp), %eax
20 cmpl -36(%rbp), %eax
21 jge .L4
22 movl -4(%rbp), %eax
23 cltq
24 #leaq 0(,%rax,4), %rdx REDUNDANT STEP IGNORED
25 leaq 0(,%rax,4), %rcx #MADE FASTER
26 movq -24(%rbp), %rax
27 addq %rcx, %rax #CHANGED IN ORDER TO USE MORE THAN ONE AT A TIME IN ORDER TO SAVE ENERGY

```

Figure 7: Assembly code for my optimized code for calculating the dot product using Indexes. In this file I have removed line 24 which is redundant because it is looking at a register to store the values more than once, that is not needed. It can continue to look at the same register and update the value that is stored there over time instead of having to continue to switch between the registers to find the total

sum. Line 25 was also optimized to make the process go even faster using the same implementation.

Line 27 was modified to make the summation process take less energy from the processor.

```

28     movl    (%rax), %edx
29     movl    -4(%rbp), %eax
30     cltq
31     #leaq   0(,%rax,4), %rcx    MOVED TO LINE 26
32     movq    -32(%rbp), %rax
33     addq    %rcx, %rax
34     movl    (%rax), %eax
35
36     imull   %edx, %eax
37     addl    %eax, -8(%rbp)
38     addl    $1, -4(%rbp)
39     jmp     .L3
40 .L4:
41     nop
42     popq    %rbp
43     .cfi_def_cfa 7, 8
44     ret
45     .cfi_endproc
46 .LFE0:
47     .size   _Z20DotProductUsingIndexPiS_i, .-_Z20DotProductUsingIndexPiS_i

```

Figure 8: My optimized assembly code for calculating the Dot product using the indexes instead of pointers for the Linux 64-bit compiler, I have moved line 31 to line 26, earlier in the program because I believe that if this condition is checked earlier the program will run quicker. Because it is checking a condition, the condition should be checked beforehand.

LINUX 64 BIT TIME VALUES FOR INDEX OPTIMIZED							
N=	10	100	1000	10000	100000	1000000	10000000
TRIAL	0.2396	0.9095	7.8694	77.1904	783.6297	5919.4531	57921.714
TRIAL	0.2509	0.8867	7.8933	77.2364	650.3334	5931.8516	57939.2392
TRIAL	0.2397	0.7516	7.9282	77.1041	682.2527	5920.5353	58186.8999
TRIAL	0.2366	0.9242	7.9068	77.1778	786.1063	5908.4519	57830.5958
TRIAL	0.2395	0.801	6.937	63.6697	706.0471	5989.7084	57984.7236
TRIAL	0.2476	0.9068	8.2007	77.1659	644.2415	5880.0702	58012.8488
TRIAL	0.2398	0.9234	7.9487	65.3068	688.7589	5921.5033	58208.8686
TRIAL	0.2484	0.9229	7.8318	74.459	766.4919	5879.3991	57834.553
TRIAL	0.2439	0.9402	7.8435	77.6407	776.0728	5917.1016	57912.7809
TRIAL	0.2346	0.8448	8.1168	77.198	734.11	5908.7739	57965.0948

Figure 9: Table that was recorded of the data when running the dot product using my optimized assembly code for Linux 64-bit compiler for the different values of N for the different sizes of the arrays. This information will later be graphed and compared to other sections of this project.

2.3 DOT PRODUCT USING POINTERS ON LINUX 64 BIT

In this section I will be computing the dot product using GCC generated assembly code on the Linux 64-bit compiler on Ubuntu using pointers. To generate the code what I need to do is use the flag -S with the GCC compiler. After it is loaded, I will link the assembly code to the main file that will be used to time how long it takes to perform the dot product using pointers on arrays of sizes; 10, 100, 1000, 10000, 100000, 1000000, and 10000000. Then I will measure and record the time it takes and store the data to later use it on Graph to compare to the assembly pointer optimized code.

```
1 void DotProductUsingPointer(int* Array1, int* Array2, int size){  
2     int i;  
3     int result = 0;  
4     for( i=0; i < size ; i+=1){  
5         result += (( *(Array1 + i)) * (*(Array2 + i)));  
6     }  
7 }
```

Figure 10: C++ code that will be used to generate the assembly code using the Linux 64-bit compiler on Ubuntu using pointers. I will then link it to the timing file that will be used, the main file that will be used to measure the time that it takes to run the Dot Product using the compiler generated assembly code using pointers. I will be doing this for different values of sized arrays.

```

1  .file "DOTPRODUCTUSINGPOINTER.cpp"
2  .text
3  .globl _Z22DotProductUsingPointerPiS_i
4  .type _Z22DotProductUsingPointerPiS_i, @function
5  _Z22DotProductUsingPointerPiS_i:
6  .LFB0:
7  .cfi_startproc
8  pushq %rbp
9  .cfi_def_cfa_offset 16
10 .cfi_offset 6, -16
11 movq %rsp, %rbp
12 .cfi_def_cfa_register 6
13 movq %rdi, -24(%rbp)
14 movq %rsi, -32(%rbp)
15 movl %edx, -36(%rbp)
16 movl $0, -4(%rbp)
17 movl $0, -8(%rbp)
18 .L3:
19 movl -8(%rbp), %eax
20 cmpl -36(%rbp), %eax
21 jge .L4
22 movl -8(%rbp), %eax
23 cltq
24 leaq 0(%rax,4), %rdx
25 movq -24(%rbp), %rax
26 addq %rdx, %rax
27 movl (%rax), %edx
28 movl -8(%rbp), %eax
29 cltq
30 leaq 0(%rax,4), %rcx
31 movq -32(%rbp), %rax
32 addq %rcx, %rax
33 movl (%rax), %eax
34 imull %edx, %eax
35 addl %eax, -4(%rbp)

```

Figure 11: Compiler generated assembly code for running dot product using pointers on Linux 64-bit compiler. We can see that the file is being compiled into assembly code from the file Dotproductusingpointer.cpp which has the instructions written in C++ that were generated into assembly.

These are the instructions that use the registers to look at the values that will be added to multiply and add into the compiler. Some of the instructions are redundant because they are repetitive and some of them are unnecessary because they make the process take even longer than they should take if the code were written more efficiently.

```

36     addl $1, -8(%rbp)
37     jmp .L3
38     .L4:
39     nop
40     popq %rbp
41     .cfi_def_cfa 7, 8
42     ret
43     .cfi_endproc
44     .LFE0:
45     .size _Z22DotProductUsingPointerPi5_i, .-_Z22DotProductUsingPointerPi5_i
46     .ident "GCC: (Ubuntu 5.4.0-6ubuntu1~16.04.4) 5.4.0 20160609"
47     .section .note.GNU-stack,"",@progbits

```

Figure 12: Compiler generated assembly code that is also used for calculating dot product using pointers on Linux 64-bit compiler. The final instructions of deallocating the memory from the stack pointer and from the stack completely and dereferencing.

LINUX 64 BIT TIME VALUES FOR POINTER							
N=	10	100	1000	10000	100000	1000000	10000000
TRIAL	0.1838	0.977	7.2893	71.5545	714.1938	6237.9745	61342.2397
TRIAL	0.2646	0.9423	7.3291	86.8798	826.8648	6241.4384	61164.0782
TRIAL	0.1901	0.9842	7.3364	81.6914	817.4812	6226.7609	61150.8543
TRIAL	0.2106	0.875	8.3267	71.5269	773.8148	6268.5479	61272.3362
TRIAL	0.2562	0.9901	8.3464	81.7703	773.6742	6247.637	61267.8586
TRIAL	0.1888	0.7296	8.3769	60.2667	681.0909	6262.3244	61258.7707
TRIAL	0.2575	0.9973	11.6487	81.6332	705.022	6216.3497	61284.4935
TRIAL	0.1976	1.0136	6.1492	71.5555	719.1089	6247.8569	61230.9624
TRIAL	0.2268	0.885	8.3284	82.3511	779.0199	6223.1534	61212.1054
TRIAL	0.243	0.8691	7.3178	60.2061	751.0914	6337.74	61279.9792

Figure 13: Table that was recorded of the data when running the dot product using the assembly code generated by Linux 64-bit compiler for the different values of N for the different sizes of the arrays for using pointers. This information will later be graphed and compared to other sections of this project.

2.4 DOT PRODUCT USING OPTIMIZED POINTER ON LINUX 64 BIT

In this section I will be computing the dot product using GCC generated assembly code on the Linux 64-bit compiler on Ubuntu that I have optimized. To generate the code what I need to do is use the flag -S with the GCC compiler. After it is loaded. I will link the assembly code to the main file that will be used to time how long it takes to perform the dot product using pointers on arrays of sizes; 10, 100, 1000, 10000, 100000, 1000000, and 10000000. Then I will measure and record the time it takes and

store the data to later use it on Graph to compare to the compiler generated assembly code for performing the dot product using pointers.

```

1 void DotProductUsingPointer(int* Array1, int* Array2, int size){
2     int i;
3     int result = 0;
4     for( i=0; i < size ; i+=1){
5         result += (( *(Array1 + i)) * (*(Array2 + i)));
6     }
7 }

```

Figure 14: C++ code that will be used to generate the assembly code using the Linux 64-bit compiler on Ubuntu and will then be optimized by me for calculating the dot product of arrays using. After I optimize the code I will then link it to the timing file that will be used, the main file that will be used to measure the time that it takes to run the Dot Product using my optimized assembly code. I will be doing this for different values of sized arrays.

```

1 .file "DOTPRODUCTUSINGPOINTER.cpp"
2 .text
3 .globl _Z22DotProductUsingPointerPiS_i
4 .type _Z22DotProductUsingPointerPiS_i, @function
5 _Z22DotProductUsingPointerPiS_i:
6 .LFB0:
7 .cfi_startproc
8 pushq %rbp
9 .cfi_def_cfa_offset 16
10 .cfi_offset 6, -16
11 movq %rsp, %rbp
12 .cfi_def_cfa_register 6
13 movq %rdi, -24(%rbp)
14 movq %rsi, -32(%rbp)
15 movl %edx, -36(%rbp)
16 movl $0, -4(%rbp)
17 movl $0, -8(%rbp)
18 .L3:
19 movl -8(%rbp), %eax
20 cmpl -36(%rbp), %eax
21 jge .L4
22 movl -8(%rbp), %eax
23 cltq
24 #leaq 0(,%rax,4), %rdx REDUNDANT INSTRUCTION
25 leaq 0(,%rax,4), %rcx #EXECUTED EARLIER
26 movq -24(%rbp), %rax
27 addq %rcx, %rax #USING A DIFFERENT REGISTER
28 movl (%rax), %edx
29 movl -8(%rbp), %eax
30 cltq

```

Figure 15: Assembly code for my optimized code for calculating the dot product using pointers. In this file I have removed line 24 which is redundant because it is looking at a register to store the values more than once, that is not needed.

It can continue to look at the same register and update the value that is stored there over time instead of having to continue to switch between the registers to find the total sum. Line 25 was also optimized to make the process go even faster using the same implementation by making it be performed earlier in the code. Line 27 was modified to make the summation process take less energy from the processor by using a different register.

```

31  #leaq 0(%rax,4), %rcx    EXECUTED ON LINE 25 INSTEAD
32  movq  -32(%rbp), %rax
33  addq  %rcx, %rax
34  movl  (%rax), %eax
35  imull %edx, %eax
36  addl  %eax, -4(%rbp)
37  addl  $1, -8(%rbp)
38  jmp  .L3
39  .L4:
40  nop
41  popq  %rbp
42  .cfi_def_cfa 7, 8
43  ret
44  .cfi_endproc
45  .LFE0:
46  .size _Z22DotProductUsingPointerPi5_i, .-_Z22DotProductUsingPointerPi5_i
47  .ident "GCC: (Ubuntu 5.4.0-6ubuntu1~14.04.4) 5.4.0 20160609"
48  .section .note.GNU-stack,"",@progbits

```

Figure 16: My optimized assembly code for calculating the Dot product using the pointers instead of pointers for the Linux 64-bit compiler, I have moved line 31 to line 27, earlier in the program because I believe that if this condition is checked earlier the program will run quicker. Because it is checking a condition, the condition should be checked beforehand which is going to make the running time significantly decrease.

LINUX 64 BIT TIME VALUES FOR POINTER OPTIMIZED							
N=	10	100	1000	10000	100000	1000000	10000000
TRIAL	0.2417	0.9369	7.9269	76.9864	748.7691	5906.7886	57699.6842
TRIAL	0.2501	0.8999	7.9221	77.0992	754.3107	5909.4796	57903.585
TRIAL	0.2392	0.9347	7.9153	77.1465	782.1948	5896.2628	57842.3739
TRIAL	0.2458	0.9158	7.92	77.6904	742.119	5914.5256	57839.8326
TRIAL	0.2402	0.9262	7.8654	77.1256	748.7056	5895.6659	57802.6829

TRIAL	0.2593	0.9065	7.8748	77.1631	787.069	5894.6824	57882.6436
TRIAL	0.2472	0.9076	7.9426	77.4981	743.5704	5906.9981	57885.6313
TRIAL	0.2419	0.8961	7.867	76.9301	736.5196	5928.273	57898.9538
TRIAL	0.2364	0.9358	7.8643	77.1463	744.3211	5924.4438	57703.2393
TRIAL	0.2446	0.9193	7.8952	77.8602	717.1894	5908.8214	57940.9472

Figure 17: Results for the times of calculating the dot product for different values of N using my optimized code and pointers on Linux 64-bit compiler.

3. DOT PRODUCT ON LINUX 32 BIT

In this section I will be calculating dot product using different methods. I will be doing it on Linux 32-bit compiler and I will be using a Raspberry Pi. I will calculate it using compiler generated assembly code, that is generated for calculating using indexes and pointer and I will optimize both to measure more time and efficiency.

```

2  #include <stdio.h>
3  #include <time.h>
4  #include <stdlib.h>
5  #include <stdint.h>
6  // #include <iostream>
7  using namespace std;
8  //void DotProductUsingIndex(int Array1[],int Array2[], int size);
9  void DotProductUsingPointer(int *Array1,int *Array2, int size);
10 //Size = 10, 100, 1000, 10000, 100000, 1000000, 10000000
11 #define SIZE 10
12 static int Array1[SIZE];
13 static int Array2[SIZE];
14 main(int argc, char **argv){
15     uint64_t sum = 0.0;
16     for (int i=0;i<SIZE;i++){Array1[i]=10;Array2[i]=15;}
17     for(int i = 0; i < 10; i++){
18         timespec start, end;
19         clock_gettime(CLOCK_MONOTONIC, &start);
20         // DotProductUsingIndex(Array1,Array2, SIZE);
21         DotProductUsingPointer(Array1,Array2, SIZE);
22         clock_gettime(CLOCK_MONOTONIC,&end);
23         sum += (end.tv_nsec - start.tv_nsec);
24     }
25     printf("size = %i, average is %f ms \n", SIZE,(double)(sum/10.0));
26     return 0;
27 }

```

Figure 18: The main file that will be used in each of the following 4 sections in order to measure the time of running each of the dot product methods on a Raspberry pi for Linux 32-bit.

3.1 DOT PRODUCT USING INDEXES ON LINUX 32 BIT.

In this section I will be computing the dot product using GCC generated assembly code on the Linux 32-bit compiler on a Raspberry Pi. To generate the code what I need to do is use the flag -S with the GCC compiler as well. After it is loaded, I will link the assembly code to the main file that will be used to time how long it takes to perform the dot product using indexes on arrays of sizes; 10, 100, 1000, 10000, 100000, 1000000, and 10000000. Then I will measure and record the time it takes and store the data to later use it on Graph to compare to its optimized code.

```

1 void DotProductUsingIndex(int A[],int B[],int size ){
2     int result = 0;
3     for(int i = 0; i < size; i++)
4         result += ( A[i] * B[i]);
5 }

```

Figure 19: C++ code that will be used in for generating the dot product assembly code that will then be linked to the main file on a Raspberry Pi.

```

1 .arch armv6
2 .eabi_attribute 27, 3
3 .eabi_attribute 28, 1
4 .fpu vfp
5 .eabi_attribute 20, 1
6 .eabi_attribute 21, 1
7 .eabi_attribute 23, 3
8 .eabi_attribute 24, 1
9 .eabi_attribute 25, 1
10 .eabi_attribute 26, 2
11 .eabi_attribute 30, 6
12 .eabi_attribute 34, 1
13 .eabi_attribute 18, 4
14 .file "DotProductUsingIndex.cpp"
15 .text
16 .align 2
17 .global _Z20DotProductUsingIndexPiS_i
18 .type _Z20DotProductUsingIndexPiS_i, %function
19 _Z20DotProductUsingIndexPiS_i:
20 .fnstart
21 .LFB0:
22 @ args = 0, pretend = 0, frame = 24
23 @ frame_needed = 1, uses_anonymous_args = 0
24 @ link register save eliminated.

```

Figure 20: Assembly code that is generated by compiler on Raspberry pi.

```

25     str fp, [sp, #-4]!
26     add fp, sp, #0
27     sub sp, sp, #28
28     str r0, [fp, #-16]
29     str r1, [fp, #-20]
30     str r2, [fp, #-24]
31     mov r3, #0
32     str r3, [fp, #-8]
33     mov r3, #0
34     str r3, [fp, #-12]
35     b .L2
36     .L3:
37     ldr r3, [fp, #-12]
38     mov r3, r3, asl #2
39     ldr r2, [fp, #-16]
40     add r3, r2, r3
41     ldr r3, [r3]
42     ldr r2, [fp, #-12]
43     mov r2, r2, asl #2
44     ldr r1, [fp, #-20]
45     add r2, r1, r2
46     ldr r2, [r2]
47     mul r3, r2, r3

```

Figure 21: Assembly code for calculating the dot product using indexes on the raspberry pi.

```

48     ldr r2, [fp, #-8]
49     add r3, r2, r3
50     str r3, [fp, #-8]
51     ldr r3, [fp, #-12]
52     add r3, r3, #1
53     str r3, [fp, #-12]
54     .L2:
55     ldr r2, [fp, #-12]
56     ldr r3, [fp, #-24]
57     cmp r2, r3
58     blt .L3
59     sub sp, fp, #0
60     @ sp needed
61     ldr fp, [sp], #4
62     bx lr
63     .cantunwind
64     .fnend
65     .size _Z20DotProductUsingIndexPiS_i, _Z20DotProductUsingIndexPiS_i
66     .ident "GCC: (Raspbian 4.9.2-10) 4.9.2"
67     .section .note.GNU-stack,"",%progbits

```

Figure 22: Continued assembly code for calculating the dot product using indexes on the Raspberry pi
these are the final steps where the program terminates and deallocates the memory.

LINUX 32 BIT TIME VALUES FOR INDEX							
N=	10	100	1000	10000	100000	1000000	10000000
TRIAL	1.5729	6.3074	52.7863	545.5257	3711.325	30968.2505	14757395.3
TRIAL	0.8697	6.2187	27.3228	532.6872	5585.355	30895.0895	14757395.3
TRIAL	1.5782	6.2395	52.7865	539.0047	3506.81	32547.6609	14757395.3
TRIAL	0.802	6.224	53.0573	538.1557	5488.095	29809.533	14757395.3

TRIAL	1.552	6.1459	26.5885	528.5774	2855.514	30802.1625	14757395.3
TRIAL	1.5574	6.2345	53.1718	262.4059	5620.345	30306.5379	14757395.3
TRIAL	1.6407	6.1927	53.1719	535.9736	2849.769	29885.809	14757395.3
TRIAL	1.5782	6.2032	53.1407	261.0623	5609.037	31123.2299	14757395.3
TRIAL	1.5468	3.1198	52.828	544.4006	2845.805	28103.9249	14757395.3
TRIAL	1.5885	6.3436	52.7916	262.5206	7241.947	29037.1169	14757395.3

Figure 23: Time measurements for calculating the dot product when the compiler generated assembly code is linked to the main file that calculates time on the raspberry pi using different values of N for the sizes of the arrays.

3.2 DOT PRODUCT USING OPTIMIZED INDEX ON LINUX 32 BIT

In this section I will be computing the dot product using GCC generated assembly code that I have optimized on the Linux 32-bit compiler on a Raspberry Pi To generate the code what I need to do is use the flag -S with the GCC compiler. After it is loaded. I will optimize it and link the assembly code to the main file that will be used to time how long it takes to perform the dot product using indexes optimized on arrays of sizes; 10, 100, 1000, 10000, 100000, 1000000, and 10000000. Then I will measure and record the time it takes and store the data to later use it on Graph to compare to its normal assembly code.

```

1  void DotProductUsingIndex(int A[],int B[],int size ){
2      int result = 0;
3      for(int i = 0; i < size; i++)
4          result += ( A[i] * B[i]);
5  }
```

Figure 24: C++ code for calculating the Dot product using indexes on the raspberry pi. This code will be used to generate assembly code with the compiler, after it is generated I will optimize it and then calculate the dot product using indexes for various sizes of arrays.

```

1  .arch armv6
2  .eabi_attribute 28, 1
3  .eabi_attribute 20, 1
4  .eabi_attribute 21, 1
5  .eabi_attribute 23, 3
6  .eabi_attribute 24, 1
7  .eabi_attribute 25, 1
8  .eabi_attribute 26, 2
9  .eabi_attribute 30, 6
10 .eabi_attribute 34, 1
11 .eabi_attribute 18, 4
12 .file "DOTPRODUCTUSINGINDEX.cpp"
13 .text
14 .align 2
15 .global _Z20DotProductUsingIndexPiS_i
16 .syntax unified
17 .arm
18 .fpu vfp
19 .type _Z20DotProductUsingIndexPiS_i, %function
20 _Z20DotProductUsingIndexPiS_i:
21 .fnstart
22 .LF80:
23 @ args = 0, pretend = 0, frame = 24
24 @ frame_needed = 1, uses_anonymous_args = 0
25 @ link register save eliminated.
26 str fp, [sp, #-4]!
27 add fp, sp, #0
28 sub sp, sp, #28
29 str r0, [fp, #-16]
30 str r1, [fp, #-20]
31 str r2, [fp, #-24]
32 mov r3, #0
33 str r3, [fp, #-8]
34 mov r4, #0 @USING A DIFFERENT REGISTER
35 str r4, [fp, #-12] @USING A DIFFERENT REGISTER

```

Figure 25: Optimized assembly code by me for calculating the Dot product using indexes on a raspberry pi. What I did on line 34 was use a different register so that I can speed up the process. Instead of using the register that was being used before to copy a temporary value over, I am copying the recursively like using an accumulator so that it computer faster. I then use the same register on line 35 to speed up even more the assembly code for calculating the dot product using indexes on a raspberry pi.

```

36 .L3:
37 @ldr r2, [fp, #-12] @REDUNDANT STEP
38 ldr r3, [fp, #-24]
39 cmp r4, r3 @USING A DIFFERENT REGISTER
40 bge .L4
41 ldr r3, [fp, #-12]
42 lsl r3, r3, #2
43 ldr r2, [fp, #-16]
44 add r3, r2, r3
45 ldr r3, [r3]
46 ldr r2, [fp, #-12]
47 lsl r2, r2, #2
48 ldr r1, [fp, #-20]
49 add r2, r1, r2
50 ldr r2, [r2]
51 mul r3, r2, r3
52 ldr r2, [fp, #-8]
53 add r3, r2, r3
54 str r3, [fp, #-8]
55 @ldr r3, [fp, #-12] @REDUNDANT
56 add r4, r4, #1 @USING THE SAME REGISTER INSTEAD OF 2 DIFFERENT ONES
57 str r4, [fp, #-12] @USING A DIFFERENT POINTER TO STORE.
58 b .L3

```

Figure 26: assembly code for calculating the dot product using indexes on the raspberry pi, I have optimized it by removing 2 redundant steps on line 37 and 55, I also used the same register to add on line 56 and a different pointer n line 57. I also used a different register on line 39.

```

59 .L4:
60 nop
61 add sp, fp, #0
62 @ sp needed
63 ldr fp, [sp], #4
64 bx lr
65 .cantunwind
66 .fncnd
67 .size _Z20DotProductUsingIndexPi5_i, .-_Z20DotProductUsingIndexPi5_i
68 .ident "GCC: (Raspbian 6.3.0-18+rpil) 6.3.0 20170516"
69 .section .note.GNU-stack,"",%progbits

```

Figure 27: The rest of the assembly code for calculating the dot product using indexes on the raspberry pi that I have optimized, I used the stack pointer instead on line 62 instead of what it was previously using because I was certain that using the stack pointer instead would have an impact on the calculation time.

LINUX 32 BIT TIME VALUES FOR INDEX OPTIMIZED							
N=	10	100	1000	10000	100000	1000000	10000000
TRIAL	1.4844	14.047	46.5884	455.6138	4782.2605	27412.4208	48552.1957
TRIAL	0.7864	5.5365	23.073	234.442	3569.0911	26126.4709	46483.8735
TRIAL	1.4843	2.7914	46.1147	455.8635	2454.7499	29549.6349	49432.5579
TRIAL	1.4947	5.5155	25.3335	234.3222	5084.4007	26893.8493	48684.3898
TRIAL	1.5937	2.8072	46.0834	455.7335	2937.1501	25171.0252	48594.6608

TRIAL	1.4582	5.4947	46.1561	234.333	4873.9586	29456.3335	53501.1047
TRIAL	1.4844	2.7814	23.2084	227.7706	3917.377	27633.7959	53589.6603
TRIAL	1.5885	5.4948	46.3543	455.4627	5760.3211	27513.3849	48210.6327
TRIAL	1.5001	5.5104	23.0833	227.7912	3785.8201	25914.5242	74211.7123
TRIAL	0.7552	5.5157	46.5777	234.3172	4962.1299	25675.1653	48831.0885

Figure 28: Data table for the measured time for calculating the dot product using indexes with my optimized linked assembly code on the Raspberry pi using 32-bit Linux and GCC.

3.3 DOT PRODUCT USING POINTERS ON LINUX 32 BIT

In this section I will be computing the dot product using GCC generated assembly code on the Linux 32-bit compiler on a Raspberry Pi using pointers. To generate the code what I need to do is use the flag -S with the GCC compiler. After it is loaded. I will link the assembly code to the main file that will be used to time how long it takes to perform the dot product using pointers on arrays of sizes; 10, 100, 1000, 10000, 100000, 1000000, and 10000000. Then I will measure and record the time it takes and store the data to later use it on Graph to compare to the assembly pointer optimized code.

```

1 void DotProductUsingPointer(int* Array1, int* Array2, int size){
2     int i;
3     int result = 0;
4     for( i=0; i < size ; i++){
5         result += (( *(Array1 + i)) * (*(Array2 + i)));
6     }
7 }
```

Figure 29: C++ code that will be used to generate the assembly code for calculating the dot product using pointers on the raspberry pi. After I generate the assembly code I will link the code to the source file that I am using as a main file to calculate the time for each run. Then I will take measurements of time to see how quickly or how long it takes to calculate the dot product for each different size, in the next section I will optimize the code.

```

1  .arch armv6
2  .eabi_attribute 27, 3
3  .eabi_attribute 28, 1
4  .fpu vfp
5  .eabi_attribute 20, 1
6  .eabi_attribute 21, 1
7  .eabi_attribute 23, 3
8  .eabi_attribute 24, 1
9  .eabi_attribute 25, 1
10 .eabi_attribute 26, 2
11 .eabi_attribute 30, 6
12 .eabi_attribute 34, 1
13 .eabi_attribute 18, 4
14 .file "DOTPRODUCTUSINGPOINTER.cpp"
15 .text
16 .align 2
17 .global _Z22DotProductUsingPointerPiS_i
18 .type _Z22DotProductUsingPointerPiS_i, %function
19 _Z22DotProductUsingPointerPiS_i:
20 .fnstart
21 .LFB0:
22 @ args = 0, pretend = 0, frame = 24
23 @ frame_needed = 1, uses_anonymous_args = 0
24 @ link register save eliminated.
25 str fp, [sp, #-4]!
26 add fp, sp, #0
27 sub sp, sp, #28
28 str r0, [fp, #-16]
29 str r1, [fp, #-20]
30 str r2, [fp, #-24]
31 mov r3, #0
32 str r3, [fp, #-12]
33 mov r3, #0
34 str r3, [fp, #-8]
35 b .L2
36 .L3:
37 ldr r3, [fp, #-8]
38 mov r3, r3, asl #2

```

Figure 30: Assembly code for running dot product using pointers on the Linux 32bit compiler. On a raspberry pi that I will later link to the main file to take time measurements.

```

39  ldr r2, [fp, #-16]
40  add r3, r2, r3
41  ldr r3, [r3]
42  ldr r2, [fp, #-8]
43  mov r2, r2, asl #2
44  ldr r1, [fp, #-20]
45  add r2, r1, r2
46  ldr r2, [r2]
47  mul r3, r2, r3
48  ldr r2, [fp, #-12]
49  add r3, r2, r3
50  str r3, [fp, #-12]
51  ldr r3, [fp, #-8]
52  add r3, r3, #1
53  str r3, [fp, #-8]
54  .L2:
55  ldr r2, [fp, #-8]
56  ldr r3, [fp, #-24]
57  cmp r2, r3
58  blt .L3
59  sub sp, fp, #0
60  @ sp needed
61  ldr fp, [sp], #4
62  bx lr
63  .cantunwind
64  .fend
65  .size _Z22DotProductUsingPointerPiS_i, -_Z22DotProductUsingPointerPiS_i
66  .ident "GCC: (Raspbian 4.9.2-10) 4.9.2"
67  .section .note.GNU-stack,"",%progbits

```

Figure 31: Continuation of the assembly code for calculating the dot product using pointers on the Raspberry pi.

LINUX 32 BIT TIME VALUES FOR POINTER							
N=	10	100	1000	10000	100000	1000000	10000000
TRIAL	1.4947	6.2032	53.2031	538.1715	5584.2298	31913.1197	12912720.85
TRIAL	1.5207	6.2187	53.203	532.4422	3639.3249	30134.6108	12912720.85
TRIAL	1.5314	3.1406	52.8385	532.1662	2809.1905	28614.5755	12912720.85
TRIAL	0.7916	6.1824	26.5313	268.6142	6507.9321	30937.4278	12912720.85
TRIAL	1.5	4.2448	52.7604	531.8125	4638.2565	30770.2404	12912720.85
TRIAL	0.7761	6.2084	26.3855	262.38	5474.0995	29561.5544	12912720.85
TRIAL	0.7864	3.0834	52.8072	531.4215	2803.2841	31763.5157	12912720.85
TRIAL	1.5626	3.099	26.3855	262.38	5535.8132	31015.8394	12912720.85
TRIAL	1.7656	6.2031	26.3802	530.1768	4402.3608	32091.953	12912720.85
TRIAL	1.5468	6.25	52.7706	530.6664	9222.9558	30712.4802	12912720.85

Figure 32: Data segment for the measurements of calculating the dot product using pointers on the Raspberry pi. Some of these measurements will be graphed and compared to the other platforms that are being tested in this project and to the optimized measurements.

3.4 DOT PRODUCT USING OPTIMIZED POINTER ON LINUX 32 BIT

In this section I will be computing the dot product using GCC generated assembly code on the Linux 32-bit compiler on Raspberry Pi that I have optimized. To generate the code what I need to do is use the flag -S with the GCC compiler. After it is loaded. I will link the assembly code to the main file that will be used to time how long it takes to perform the dot product using pointers on arrays of sizes; 10, 100, 1000, 10000, 100000, 1000000, and 10000000. Then I will measure and record the time it takes and store the data to later use it on Graph to compare to the compiler generated assembly code for performing the dot product using pointers.

```
1 void DotProductUsingPointer(int* Array1, int* Array2, int size){  
2     int i;  
3     int result = 0;  
4     for( i=0; i < size ; i+=1){  
5         result += (( *(Array1 + i)) * (*(Array2 + i)));  
6     }  
7 }
```

Figure 33: C++ code that will be used to generate the assembly code for calculating the dot product of n size arrays, that I will then optimize and link to the main file that will be used to measure the time that it takes. The original assembly code that will be optimized by me is initially generated by the compiler that is 32 bits on Linux on the raspberry pi, but I will then optimize it.

```

1  .arch armv6
2  .eabi attribute 28, 1
3  .eabi attribute 20, 1
4  .eabi attribute 21, 1
5  .eabi attribute 23, 3
6  .eabi attribute 24, 1
7  .eabi attribute 25, 1
8  .eabi attribute 26, 2
9  .eabi attribute 30, 6
10 .eabi attribute 34, 1
11 .eabi attribute 18, 4
12 .file "DotProductUsingPointer.cpp"
13 .text
14 .align 2
15 .global _Z22DotProductUsingPointerPiS_i
16 .syntax unified
17 .arm
18 .fpu vfp
19 .type _Z22DotProductUsingPointerPiS_i, %function
20 _Z22DotProductUsingPointerPiS_i:
21 .fnstart
22 .LFB0:
23 @ args = 0, pretend = 0, frame = 24
24 @ frame needed = 1, uses anonymous args = 0
25 @ link register save eliminated.
26 str fp, [sp, #-4]!
27 add fp, sp, #0
28 sub sp, sp, #28
29 str r0, [fp, #-16]
30 str r1, [fp, #-20]
31 str r2, [fp, #-24]
32 mov r3, #0
33 str r3, [fp, #-12]
34 mov r4, #0 @USING A DIFFERENT REGISTER
35 str r4, [fp, #-8] @USING A DIFFERENT REGISTER
36 .L3:
37 @ldr r2, [fp, #-8] @REDUNDANT STEP
38 ldr r3, [fp, #-24]

```

Figure 34: Assembly code for calculating the dot product using pointers that I have optimized. In line 34 I have used a different register like in the previous example. In the line 35 I have also used a different register and on line 37 I have removed the unnecessary step of using register r2 again, it is not needed because I am using a different register in the previous instruction.


```

39     cmp r4, r3
40     bge .L4
41     ldr r3, [fp, #-8]
42     lsl r3, r3, #2
43     ldr r2, [fp, #-16]
44     add r3, r2, r3
45     ldr r3, [r3]
46     ldr r2, [fp, #-8]
47     lsl r2, r2, #2
48     ldr r1, [fp, #-20]
49     add r2, r1, r2
50     ldr r2, [r2]
51     mul r3, r2, r3
52     ldr r2, [fp, #-12]
53     add r3, r2, r3
54     str r3, [fp, #-12]
55     ldr r3, [fp, #-8]    @CHANGED REGISTER
56     add r4, r4, #1
57     str r4, [fp, #-8]
58     b .L3
59
60 .L4:
61     nop
62     add sp, fp, #0
63     @ sp needed
64     ldr fp, [sp], #4
65     bx lr
66     .cantunwind
67     .fnend
68     .size _Z22DotProductUsingPointerPiS_i, .-_Z22DotProductUsingPointerPiS_i
69     .ident "GCC: (Raspbian 6.3.0-18+rpil) 6.3.0 20170516"
70     .section .note.GNU-stack,"",%progbits

```

Figure 35: Continued assembly code that I optimized for calculating the dot product using pointers on the raspberry pi. In line 55 I have changed the register that is being used so that instead of storing the value into a temporary register it uses the same register and saves time.

LINUX 32 BIT TIME VALUES FOR POINTER OPTIMIZED							
N=	10	100	1000	10000	100000	1000000	10000000
TRIAL	1.4739	5.6823	47.8072	486.8433	5907.8336	28070.5658	1660206.967
TRIAL	1.5314	5.7032	23.9062	478.0361	2593.5761	25804.2758	1660206.967
TRIAL	0.7501	2.8594	47.8126	236.1719	5047.7824	28059.0967	1660206.967
TRIAL	1.5105	5.6875	24.1093	477.4583	2554.9668	29735.7571	1660206.967
TRIAL	0.7344	5.8074	47.7604	236.25	7137.5308	30443.996	1660206.967
TRIAL	1.5001	5.7032	23.9688	494.7183	2589.9927	26019.6767	1660206.967
TRIAL	0.7187	2.8699	47.7552	237.3956	5042.9435	25423.2865	1660206.967
TRIAL	1.5417	5.6613	47.8645	490.859	2572.8832	27562.9879	1660206.967
TRIAL	1.5	2.9895	47.8855	237.0831	5061.121	29667.2624	1660206.967
TRIAL	1.4373	5.6719	23.8801	481.5622	2572.4666	28223.2373	1660206.967

Figure 36: Data table for the measurements of calculating the dot product on a Linux 32-bit compiler using my optimized assembly code for different values of N or the different values for the sizes of the arrays.

4. DOT PRODUCT ON WINDOWS 32 BIT

In this section I will be calculating the dot product and measure the time that it takes to run the dot product on windows 32-bit compiler. I will test it using indexes and pointers, linking using assembly code and I will use the code that is generated but I will also use the assembly code that I will optimize.

In the final section I will graph and compare the data to show that my optimizations were effective.

```

1  #include "stdafx.h"
2  #include <tchar.h>
3  #include <windows.h>
4  #include <iostream>
5  // int DotProductIndex(int[], int[], int n);
6  void DotProductPointer(int *, int *, int n);
7  const int n = 10000000;
8  static int x[n];
9  static int y[n];
10 using namespace std;
11 int main()
12 {
13     __int64 ctr1 = 0, ctr2 = 0, freq = 0;
14     int acc = 0, i = 0;
15     double time = 0.0;
16     int sum;
17     int j;
18     for (i = 0; i < 100; i++)
19     {
20         for (j = 0; j < n; j++) {
21             x[j] = 2;
22             y[j] = 3;
23         }
24         if (QueryPerformanceCounter((LARGE_INTEGER *)&ctr1) != 0)

```

Figure 37: C++ code for measuring the time that it will take to calculate the dot product for each of the 4 incoming sections.

```

25     {
26         // sum = DotProductIndex(x, y, n);
27         DotProductPointer(x, y, n);
28         QueryPerformanceCounter((LARGE_INTEGER *)&ctr2);
29         cout << "Start Value: {0}" << ctr1 << endl;
30         cout << "End Value: {0}" << ctr2 << endl;
31         QueryPerformanceFrequency((LARGE_INTEGER *)&freq);
32         time += (((ctr2 - ctr1)*1.0 / freq));
33         cout << "QueryPerformanceCounter minimum resolution: 1/{0} Seconds." << freq << endl;
34         cout << n << " Increment time: {0} seconds." << (ctr2 - ctr1) * 1.0 / freq << endl; // changed size to n
35         cout << "End Value - Start Value = " << ctr2 - ctr1 << endl;
36     }

```

Figure 38: Continued C++ code for measuring time on the windows 32-bit compiler.

```

37     else
38     {
39         DWORD dwError = GetLastError();
40         cout << "Error value = {0}" << dwError << endl;
41     }
42 }
43 time = time / 100;
44 cout << "Average time is: " << time << endl;
45 system("PAUSE");
46 return 0;
47 }

```

Figure 39: Continued C++ code for measuring the time on windows 32-bit compiler for dot products.

4.1 DOT PRODUCT USING INDEXES ON WINDOWS 32 BIT.

In this section I will be computing the dot product using Visual Studio and Windows 32-bit compiler generated assembly code using Indexes. To generate the code what I need to do edit the properties of the Compiler to generate assembly code. This will happen by selecting properties and then selecting output assembler FA to see the assembly instructions that will be stored into the debug folder. After it is loaded, I will link the assembly code to the main file by using the Windows Macro Assembler setting. Then it will be used to time how long it takes to perform the dot product using indexes on arrays of sizes; 10, 100, 1000, 10000, 100000, 1000000, and 10000000. Then I will measure and record the time it takes and store the data to later use it on Graph to compare to its optimized code.

```

1 void DotProductIndex(int ArrayA[], int ArrayB[], const int n)
2 {
3     int i;
4     int sum = 0;
5     for (i = 0; i < n; i++)
6     {
7         sum += ArrayA[i] * ArrayB[i];
8     }
9 }

```

Figure 40: C++ code that will be used to generate the assembly code for running the dot product using indexes on Windows 32-bit compiler using visual studio.

```

1 ; Listing generated by Microsoft (R) Optimizing Compiler Version 19.11.25507.1
2 TITLE C:\Users\Jeter_Gutierrez\Documents\DotProductIndex.cpp
3 .686P
4 .XMM
5 include listing.inc
6 .model flat
7 INCLUDELIB MSVCRTD
8 INCLUDELIB OLDNAMES
9 PUBLIC ?DotProductIndex@@YAXQAHOH@Z ; DotProductIndex
10 EXTRN __RTC_InitBase:PROC
11 EXTRN __RTC_Shutdown:PROC
12 ; COMDAT rtc$TMZ
13 rtc$TMZ SEGMENT
14 __RTC_Shutdown.rtc$TMZ DD FLAT: __RTC_Shutdown
15 rtc$TMZ ENDS
16 ; COMDAT rtc$IMZ
17 rtc$IMZ SEGMENT
18 __RTC_InitBase.rtc$IMZ DD FLAT: __RTC_InitBase
19 rtc$IMZ ENDS
20 ; Function compile flags: /Odtp /RTCsu /ZI
21 ; COMDAT ?DotProductIndex@@YAXQAHOH@Z
22 _TEXT SEGMENT
23 _sum$ = -20 ; size = 4
24 _i$ = -8 ; size = 4
25 _ArrayA$ = 8 ; size = 4
26 _ArrayB$ = 12 ; size = 4
27 _n$ = 16 ; size = 4
28 ?DotProductIndex@@YAXQAHOH@Z PROC ; DotProductIndex, COMDAT
29 ; File C:\Users\Jeter_Gutierrez\Documents\DotProductIndex.cpp
30 ; Line 6
31 push ebp

```

Figure 41: Compiler generated assembly code for running the dot product using indexes on the Windows 32-bit compiler using visual studio.

```

31  push  ebp
32  mov  ebp, esp
33  sub  esp, 216      ; 000000d8H
34  push  ebx
35  push  esi
36  push  edi
37  lea  edi, DWORD PTR [ebp-216]
38  mov  ecx, 54      ; 00000036H
39  mov  eax, -858993460 ; cccccccc
40  rep  stosd
41  ; Line 8
42  mov  DWORD PTR _sum$[ebp], 0
43  ; Line 9
44  mov  DWORD PTR _i$[ebp], 0
45  jmp  SHORT $LN4@DotProduct
46  $LN2@DotProduct:
47  mov  eax, DWORD PTR _i$[ebp]
48  add  eax, 1
49  mov  DWORD PTR _i$[ebp], eax
50  $LN4@DotProduct:
51  mov  eax, DWORD PTR _i$[ebp]
52  cmp  eax, DWORD PTR _n$[ebp]
53  jge  SHORT $LN1@DotProduct
54  ; Line 11
55  mov  eax, DWORD PTR _i$[ebp]
56  mov  ecx, DWORD PTR _ArrayA$[ebp]
57  mov  edx, DWORD PTR _i$[ebp]
58  mov  esi, DWORD PTR _ArrayB$[ebp]
59  mov  eax, DWORD PTR [ecx+eax*4]
60  imul eax, DWORD PTR [esi+edx*4]

```

Figure 42: Continued assembly code generated by the visual studio windows 32-bit compiler for calculating the dot product using indexes on the.

```

61  add  eax, DWORD PTR _sum$[ebp]
62  mov  DWORD PTR _sum$[ebp], eax
63  ; Line 12
64  jmp  SHORT $LN2@DotProduct
65  $LN1@DotProduct:
66  ; Line 14
67  pop  edi
68  pop  esi
69  pop  ebx
70  mov  esp, ebp
71  pop  ebp
72  ret  0
73  ?DotProductIndex@@YAXQAHH@Z ENDP ; DotProductIndex
74  _TEXT ENDS
75  END

```

Figure 43: Continued assembly code generated by the visual studio windows 32-bit compiler for calculating the dot product using indexes on the that is now finalized and will be linked to the main file

to measure the time of running the dot product using indexes on windows 32-bit compile for various values of N.

WINDOWS 32 BIT TIME VALUES FOR INDEX							
N=	10	100	1000	10000	100000	1000000	10000000
TRIAL	1	1	8	79	1051	6512	67008
TRIAL	1	1	8	80	1033	6655	67509
TRIAL	1	1	8	80	1188	6538	71105
TRIAL	1	1	8	79	1034	9152	71239
TRIAL	1	1	8	77	1032	6882	66201
TRIAL	1	1	8	79	1131	6610	67026
TRIAL	1	1	8	81	1076	6680	69400
TRIAL	1	1	8	80	1052	6940	65863
TRIAL	1	1	8	79	1124	6521	66176
TRIAL	1	1	8	81	1033	6680	67351

Figure 44: Measured time data for calculating the dot product for different values of N on the Windows 32-bit compiler using visual studio. This data will later be graphed and compared to the optimized code and to the other platforms.

4.2 DOT PRODUCT USING OPTIMIZED INDEX ON WINDOWS 32 BIT

In this section I will be computing the dot product using Visual Studio and Windows 32-bit compiler generated assembly code using Indexes. To generate the code what I need to do edit the properties of the Compiler to generate assembly code which I will then optimize. This will happen by selecting properties and then selecting output assembler FA to see the assembly instructions that will be stored into the debug folder. After it is loaded, I will link the assembly code to the main file by using the Windows Macro Assembler setting. Then it will be used to time how long it takes to perform the dot product using indexes on arrays of sizes; 10, 100, 1000, 10000, 100000, 1000000, and 10000000. Then I will measure and record the time it takes and store the data to later use it on Graph to compare to its compiler generated assembly code.

```

1 void DotProductIndex(int ArrayA[], int ArrayB[], const int n)
2 {
3     int i;
4     int sum = 0;
5     for (i = 0; i < n; i++)
6     {
7         sum += ArrayA[i] * ArrayB[i];
8     }
9 }

```

Figure 45: C++ code that will be used to generate the assembly code that will be optimized by me, linked to the main file and used to measure the time it takes to calculate the dot product using indexes with my optimized assembly code.

```

1 ; Listing generated by Microsoft (R) Optimizing Compiler Version 19.00.23918.0
2 TITLE C:\Users\Jeter_Gutierrez\Documents\DotProductIndex.cpp
3 .686P
4 .XMM
5 include listing.inc
6 .model flat
7 INCLUDELIB MSVCRTD
8 INCLUDELIB OLDNAMES
9 PUBLIC ?DotProductIndex@@YAXQAHH@Z ; DotProductIndex
10 EXTRN __RTC_InitBase:PROC
11 EXTRN __RTC_Shutdown:PROC
12 ; COMDAT rtc$TMZ
13 ; rtc$TMZ SEGMENT ; IGNORING REDUNDANT STEP
14 ; __RTC_Shutdown.rtc$TMZ DD FLAT: __RTC_Shutdown ; IGNORING REDUNDANT STEP
15 ; rtc$TMZ ENDS ; IGNORING REDUNDANT STEP
16 ; COMDAT rtc$IMZ
17 ; rtc$IMZ SEGMENT ; IGNORING REDUNDANT STEP
18 ; __RTC_InitBase.rtc$IMZ DD FLAT: __RTC_InitBase ; IGNORING REDUNDANT STEP
19 ; rtc$IMZ ENDS ; IGNORING REDUNDANT STEP
20 ; Function compile flags: /Odtp /RTCsu /ZI
21 ; COMDAT ?DotProductIndex@@YAXQAHH@Z
22 _TEXT SEGMENT
23 _sum$ = -20 ; size = 4
24 _i$ = -8 ; size = 4
25 _ArrayA$ = 8 ; size = 4
26 _ArrayB$ = 12 ; size = 4
27 _n$ = 16 ; size = 4
28 ?DotProductIndex@@YAXQAHH@Z PROC ; DotProductIndex, COMDAT
29 ; File C:\Users\Jeter_Gutierrez\Documents\DotProductIndex.cpp
30 ; Line 4

```

Figure 46: Assembly code that is optimized by me, lines 13-19 are redundant I don't need them they do not affect the process of us running the code, I removed the steps that were repetitive and redundant on the visual studio assembly code.

```

31  push  ebp
32  mov  ebp, esp
33  sub  esp, 216      ; 000000d8H
34  push  ebx
35  push  esi
36  push  edi
37  lea  edi, DWORD PTR [ebp-216]
38  mov  ecx, 54      ; 00000036H
39  mov  eax, -858993460 ; ccccccccH
40  rep  stosd
41  ; Line 6
42  mov  DWORD PTR _sum$[ebp], 0
43  ; Line 7
44  ; mov DWORD PTR _i$[ebp], 0      ; IGNORING REDUNDANT STEP
45  mov  eax, 0          ; STORING I IN EAX
46  mov  edx, DWORD PTR _n$[ebp]    ; STORE SIZE INTO THE STACK
47  jmp  SHORT $LN4@DotProduct
48  $LN2@DotProduct:
49  ; mov eax, DWORD PTR _i$[ebp]    ; IGNORING REDUNDANT STEP
50  add  eax, 1
51  ; mov DWORD PTR _i$[ebp], eax
52  $LN4@DotProduct:
53  ; mov eax, DWORD PTR _i$[ebp]    ; IGNORING REDUNDANT STEP
54  ; cmp eax, DWORD PTR _n$[ebp]    ; IGNORING REDUNDANT STEP and replace with:
55  cmp  eax, edx
56  jge  SHORT $LN1@DotProduct
57  ; Line 9

```

Figure 47: The continued assembly code that I have optimized for calculating the dot product using indexes. I have removed the steps that are redundant, there were several repetitive steps that were happening, I ignored them. There were other values that were stored into different registers than the ones that the compiler had determined them to be because I determined it would be quicker to store them into eax instead of where they were being stored before.


```

58 ; mov eax, DWORD PTR _i$[ebp] ; IGNORING REDUNDANT STEP
59 mov ecx, DWORD PTR _ArrayA$[ebp]
60 ; mov edx, DWORD PTR _i$[ebp] ; IGNORING REDUNDANT STEP
61 ; mov esi, DWORD PTR _ArrayB$[ebp] ; IGNORING REDUNDANT STEP
62 ; mov eax, DWORD PTR [ecx+eax*4] ; IGNORING REDUNDANT STEP
63 mov esi, DWORD PTR [ecx+eax*4] ; THE FIRST ARRAY
64 mov ecx, DWORD PTR _Array8$[ebp] ; USING ECX INSTEAD A DIFFERENT REGISTER
65 ; imul eax, DWORD PTR [esi+edx*4]
66 imul esi, DWORD PTR [ecx+eax*4]
67 ; add eax, DWORD PTR _sum$[ebp]
68 add esi, DWORD PTR _sum$[ebp]
69 ; mov DWORD PTR _sum$[ebp], eax
70 mov DWORD PTR _sum$[ebp], esi
71 ; Line 10
72 jmp SHORT $LN2@DotProduct
73 $LN1@DotProduct:
74 ; Line 12
75 pop edi
76 pop esi
77 pop ebx
78 mov esp, ebp
79 pop ebp
80 ret 0
81 ?DotProductIndex@@YAXQAHH@Z ENDP ; DotProductIndex
82 _TEXT ENDS
83 END

```

Figure 48: Assembly code for calculating the dot product using indexes that I have optimized. I have removed more redundant steps and used a different register in line 64. This will make the program run faster.

WINDOWS 32 BIT TIME VALUES FOR INDEX OPTIMIZED							
N=	10	100	1000	10000	100000	1000000	10000000
TRIAL	1	1	9	82	800	9174	50000
TRIAL	1	1	8	88	850	5000	51324
TRIAL	1	1	10	90	900	5345	54312
TRIAL	1	1	8	100	937	5565	59872
TRIAL	1	1	6	120	956	5454	53211
TRIAL	1	1	9	127	820	5134	51234
TRIAL	1	1	8	130	825	5222	58792
TRIAL	1	1	8	124	872	5321	58797
TRIAL	1	1	6	78	778	5768	58972
TRIAL	1	1	7	80	801	5673	54312

Figure 49: measured data after I had run the dot product using my optimized code for windows 32-bit compiler, the data will be graphed later in this report and compared to the original time measurements for using the compiler generated assembly code that I have already linked and measured.

This will make it easier for us to notice that my optimized assembly code is efficient, more efficient than the code that was generated by the windows 32-bit compiler, but this is an improvement, it is more efficient.

4.3 DOT PRODUCT USING POINTERS ON WINDOWS 32 BIT

In this section I will be computing the dot product using Visual Studio and Windows 32-bit compiler generated assembly code using pointers. To generate the code what I need to do edit the properties of the Compiler to generate assembly code. This will happen by selecting properties and then selecting output assembler FA to see the assembly instructions that will be stored into the debug folder. After it is loaded, I will link the assembly code to the main file by using the Windows Macro Assembler setting. Then it will be used to time how long it takes to perform the dot product using pointers on arrays of sizes; 10, 100, 1000, 10000, 100000, 1000000, and 10000000. Then I will measure and record the time it takes and store the data to later use it on Graph to compare to its optimized code.

```
1 void DotProductPointer(int* ArrayA, int* ArrayB, const int n)
2 {
3     int i;
4     int* sum = 0;
5     for (i = 0; i < n; i++)
6     {
7         sum += *(ArrayA + i) * *(ArrayB + i);
8     }
9 }
```

Figure 50: C++ code that will be used to generate the assembly code that will be linked to the main file to measure the time it takes to calculate the dot product on windows 32-bit compiler using pointers.

```

1 ; Listing generated by Microsoft (R) Optimizing Compiler Version 19.11.25
2 TITLE C:\Users\Jeter_Gutierrez\Documents\DotProductPOINTER.cpp
3 .686P
4 .XMM
5 include listing.inc
6 .model flat
7 INCLUDELIB MSVCRTD
8 INCLUDELIB OLDNAMES
9 PUBLIC ?DotProductPointer@@YAXPAH@Z ; DotProductPointer
10 EXTRN __RTC_InitBase:PROC
11 EXTRN __RTC_Shutdown:PROC
12 ; COMDAT rtc$TMZ
13 rtc$TMZ SEGMENT
14 ; __RTC_Shutdown.rtc$TMZ DD FLAT: __RTC_Shutdown
15 rtc$TMZ ENDS
16 ; COMDAT rtc$IMZ
17 rtc$IMZ SEGMENT
18 ; __RTC_InitBase.rtc$IMZ DD FLAT: __RTC_InitBase
19 rtc$IMZ ENDS
20 ; Function compile flags: /Odtp /RTCsu /ZI
21 ; COMDAT ?DotProductPointer@@YAXPAH@Z
22 _TEXT SEGMENT
23 _sum$ = -20 ; size = 4
24 _i$ = -8 ; size = 4
25 _ArrayA$ = 8 ; size = 4
26 _ArrayB$ = 12 ; size = 4
27 _n$ = 16 ; size = 4
28 ?DotProductPointer@@YAXPAH@Z PROC ; DotProductPointer, COMDAT
29 ; File C:\Users\Jeter_Gutierrez\Documents\DotProductPOINTER.cpp
30 ; Line 6

```

Figure 51: Assembly code that was generated by the visual studio compiler on windows 32-bit compiler to calculate the dot product using pointers that will then be linked and measured.

```

31 push ebp
32 mov ebp, esp
33 sub esp, 216 ; 000000d8H
34 push ebx
35 push esi
36 push edi
37 lea edi, DWORD PTR [ebp-216]
38 mov ecx, 54 ; 00000036H
39 mov eax, -858993460 ; ccccccccH
40 rep stosd
41 ; Line 8
42 mov DWORD PTR _sum$[ebp], 0
43 ; Line 9
44 mov DWORD PTR _i$[ebp], 0
45 jmp SHORT $LN4@DotProduct
46 $LN2@DotProduct:
47 mov eax, DWORD PTR _i$[ebp]
48 add eax, 1
49 mov DWORD PTR _i$[ebp], eax
50 $LN4@DotProduct:
51 ;mov eax, DWORD PTR _i$[ebp] ;remove
52 cmp eax, DWORD PTR _n$[ebp]
53 jge SHORT $LN1@DotProduct
54 ; Line 11
55 mov eax, DWORD PTR _i$[ebp]
56 mov ecx, DWORD PTR _ArrayA$[ebp]
57 mov edx, DWORD PTR _i$[ebp]
58 mov esi, DWORD PTR _ArrayB$[ebp]
59 mov eax, DWORD PTR [ecx+eax*4]
60 imul eax, DWORD PTR [esi+edx*4]
61 mov ecx, DWORD PTR _sum$[ebp]
62 lea edx, DWORD PTR [ecx+eax*4]
63 mov DWORD PTR _sum$[ebp], edx
64 ; Line 12

```

Figure 52: Continued assembly code for generating dot product using pointers on windows 32-bit compiler.

```

65     jmp SHORT $LN2@DotProduct
66     $LN1@DotProduct:
67     ; Line 14
68     pop edi
69     pop esi
70     pop ebx
71     mov esp, ebp
72     pop ebp
73     ret 0
74     ?DotProductPointer@@YAXPAH0H@Z ENDP ; DotProductPointer
75     _TEXT ENDS
76     END

```

Figure 53: Continued assembly code for generating dot product using pointers on windows 32-bit compiler that is now finalized and will be tested.

WINDOWS 32 BIT TIME VALUES FOR POINTER							
N=	10	100	1000	10000	100000	1000000	10000000
TRIAL	1	1	8	80	808	6601	65981
TRIAL	1	1	8	80	803	6679	98933
TRIAL	1	1	8	80	850	7101	66392
TRIAL	1	1	8	80	946	6602	68729
TRIAL	1	1	8	80	815	6983	67558
TRIAL	1	1	8	80	817	6737	94812
TRIAL	1	1	8	80	806	6602	81446
TRIAL	1	1	8	80	803	7101	68279
TRIAL	1	1	8	80	969	6601	69072
TRIAL	1	1	8	80	949	6595	68279

Figure 54: Data for calculating the dot product using pointers using the compiler generated assembly code that was linked to the main file to measure the time that it takes to run the operation for different values for the sizes of the arrays.

4.4 DOT PRODUCT USING OPTIMIZED POINTER ON WINDOWS 32 BIT

In this section I will be computing the dot product using Visual Studio and Windows 32-bit compiler generated assembly code using pointers. To generate the code what I need to do edit the properties of the Compiler to generate assembly code which I will then optimize. This will happen by selecting properties and then selecting output assembler FA to see the assembly instructions that will be stored

into the debug folder. After it is loaded, I will link the assembly code to the main file by using the Windows Macro Assembler setting. Then it will be used to time how long it takes to perform the dot product using pointers on arrays of sizes; 10, 100, 1000, 10000, 100000, 1000000, and 10000000. Then I will measure and record the time it takes and store the data to later use it on Graph to compare to its compiler generated assembly code.

```

1 void DotProductPointer(int* ArrayA, int* ArrayB, const int n)
2 {
3     int i;
4     int* sum = 0;
5     for (i = 0; i < n; i++)
6     {
7         sum += *(ArrayA + i) * *(ArrayB + i);
8     }
9 }

```

Figure 55: C++ code that will be used to generate the assembly code that I will optimize for calculating the dot product using pointer on windows 32-bit compiler using visual studio.

```

1 ; Listing generated by Microsoft (R) Optimizing Compiler Version 19.00.23918.0
2 TITLE C:\Users\Jeter_Gutierrez\Documents\DotProductPOINTER.cpp
3 .XMM
4 include listing.inc
5 .model flat
6 INCLUDLIB MSVCRTD
7 INCLUDLIB OLDNAMES
8 PUBLIC ?DotProductPointer@@YAXPAH0H@Z ; DotProductPointer
9 EXTRN __RTC_InitBase:PROC
10 EXTRN __RTC_Shutdown:PROC
11 ; COMDAT rtc$TMZ
12 ; rtc$TMZ SEGMENT ; IGNORING REDUNDANT STEP
13 ; __RTC_Shutdown.rtc$TMZ DD FLAT: __RTC_Shutdown ; IGNORING REDUNDANT STEP
14 ; rtc$TMZ ENDS ; IGNORING REDUNDANT STEP
15 ; COMDAT rtc$IMZ
16 ; rtc$IMZ SEGMENT ; IGNORING REDUNDANT STEP
17 ; __RTC_InitBase.rtc$IMZ DD FLAT: __RTC_InitBase ; IGNORING REDUNDANT STEP
18 ; rtc$IMZ ENDS ; IGNORING REDUNDANT STEP
19 ; Function compile flags: /Odtp /RTCsu /ZI
20 ; COMDAT ?DotProductPointer@@YAXPAH0H@Z
21 _TEXT SEGMENT
22 _sum$ = -20 ; size = 4
23 _i$ = -8 ; size = 4
24 _ArrayA$ = 8 ; size = 4
25 _ArrayB$ = 12 ; size = 4
26 _n$ = 16 ; size = 4
27 ?DotProductPointer@@YAXPAH0H@Z PROC ; DotProductPointer, COMDAT
28 ; File C:\Users\Jeter_Gutierrez\Documents\DotProductPOINTER.cpp
29 ; Line 5
30 push ebp

```

Figure 56: Assembly code where I optimized calculating the dot product on windows 32-bit compiler using pointers. I have removed the repeated steps that are not needed to calculate the dot product.

```

31  mov ebp, esp
32  sub esp, 216          ; 000000d8H
33  push ebx
34  push esi
35  push edi
36  lea edi, DWORD PTR [ebp-216]
37  mov ecx, 54          ; 00000036H
38  mov eax, -858993460   ; ccccccccH
39  rep stosd
40  ; Line 7
41  mov DWORD PTR _sum$[ebp], 0
42  ; Line 8
43  ; mov DWORD PTR _i$[ebp], 0          ; IGNORING REDUNDANT STEP
44  mov eax, 0             ; START THE FOR LOOP
45  mov edx, DWORD PTR _n$[ebp]      ; STORE THE SIZE INTO EDX
46  jmp SHORT $LN4@DotProduct
47  $LN2@DotProduct:
48  ; mov eax, DWORD PTR _i$[ebp]      ; IGNORING REDUNDANT STEP
49  add eax, 1
50  ; mov DWORD PTR _i$[ebp], eax      ; IGNORING REDUNDANT STEP
51  $LN4@DotProduct:
52  ; mov eax, DWORD PTR _i$[ebp]      ; IGNORING REDUNDANT STEP
53  ; cmp eax, DWORD PTR _n$[ebp]      ; IGNORING REDUNDANT STEP and replace:
54  cmp eax, edx
55  jge SHORT $LN1@DotProduct
56  ; Line 10
57  ; mov eax, DWORD PTR _i$[ebp]      ; IGNORING REDUNDANT STEP
58  ; mov ecx, DWORD PTR _ArrayA$[ebp] ; IGNORING REDUNDANT STEP, START WITH SECOND ARRAY:
59  mov ecx, DWORD PTR _ArrayB$[ebp]
60  ; mov edx, DWORD PTR _i$[ebp]      ; IGNORING REDUNDANT STEP
61  ; mov esi, DWORD PTR _ArrayB$[ebp] ; IGNORING REDUNDANT STEP
62  ; mov eax, DWORD PTR [ecx+eax*4]    ; IGNORING REDUNDANT STEP
63  mov esi, DWORD PTR [ecx+eax*4]      ; SECOND ARRAY
64  mov ecx, DWORD PTR _ArrayA$[ebp]    ; SPEED UP BY UISNG ECX
65  ; imul eax, DWORD PTR [esi+edx*4]   ; IGNORING REDUNDANT STEP
66  imul esi, DWORD PTR [ecx+eax*4]      ; STORE VALUES TO ESI

```

Figure 57: I have removed most repeated steps, they do not need to happen as frequently as they did initially to calculate the dot product of arrays using pointers.

```

67  mov ecx, DWORD PTR _sum$[ebp]
68  ; lea edx, DWORD PTR [ecx+eax*4]    ; IGNORING REDUNDANT STEP and replace
69  lea edi, DWORD PTR [ecx+esi*4]
70  ; mov DWORD PTR _sum$[ebp], edx     ; IGNORING REDUNDANT STEP and replace
71  mov DWORD PTR _sum$[ebp], edi
72  ; Line 11
73  jmp SHORT $LN2@DotProduct
74  $LN1@DotProduct:
75  ; Line 13
76  pop edi
77  pop esi
78  pop ebx
79  mov esp, ebp
80  pop ebp
81  ret 0
82  ?DotProductPointer@@YAXPAH@H@Z ENDP ; DotProductPointer
83  _TEXT ENDS
84  END

```

Figure 58: Assembly code for calculating the dot product using pointers that I have optimized for the windows 32-bit compiler. I have removed 2 more redundant steps that repeat a lot.

WINDOWS 32 BIT TIME VALUES FOR POINTER OPTIMIZED							
N=	10	100	1000	10000	100000	1000000	10000000
TRIAL	1	1	6	9	100	619	6402.97
TRIAL	1	1	5	10	86	760	6478.63
TRIAL	1	1	6	8	98	820	6887.97
TRIAL	1	1	6	8	117	600	6403.94
TRIAL	1	1	5	9	78	1450	6773.51
TRIAL	1	1	6	10	79	686	6534.89
TRIAL	1	1	6	10	94	600	6403.94
TRIAL	1	1	6	7	94	627	6887.97
TRIAL	1	1	5	9	78	623	6402.97
TRIAL	1	1	5	8	79	730	6397.15

Figure 59: Time values recorded for calculating the dot product using my optimized assembly code that was optimized using pointers on Windows 32-bit compiler. This will be graphed and compared to its running time for the original code and to the other 2 platforms that were tested during this project.

5. GRAPHS

In this section I will display some graphs for the Index calculated dot product on all 3 platforms, and optimized, and using pointers and optimized for that too. We would have too many graphs if we graphed every column so instead I will only be showing the graphs for N=10, and 10000 in each case.

5.1 INDEX FOR DIFFERENT VALUES NORMAL AND OPTIMIZED ON ALL 3 PLATFORMS

In this section I will be graphing the different tables that were used to measure the time that each of the dot product performances take. I will be observing




Figure 60: Time for index on all 3 platforms for N=10;




Figure 61: Time for optimized index on all 3 platforms for N=10.

For a small value of 10 the difference in time based on the optimization is not very noticeable but it exist.




Figure 62: Time for index on all 3 platforms when N=10000.




Figure 63: Graph for time for running index optimized dot product on all 3 platforms for $n=1000$. Now using a larger value we can see that there is significant difference in the performance of the platforms for higher values of calculating the dot product using indexes, this means that I was able to successfully optimize the assembly code that was linked to the main file in every case to measure the time.

5.2 POINTER FOR DIFFERENT VALUES NORMAL AND OPTIMIZED ON ALL 3 PLATFORMS




Figure 64: Graph for the times for all 3 platforms for calculating the dot product using pointers and the compiler generated assembly code that was linked to a main file to calculate the time in each case.




Figure 65: Graph for the time that was measured when calculating the dot product on all 3 platforms using my optimized assembly code. The difference for such a small value of N where it is 10 can be seen for the raspberry pi in the fluctuation, overall we can see a difference.




Figure 66: Graph for the compiler generated assembly code that I made on all three platforms. This is for N= 10000.




Figure 67: As we can see when N is equal to 10000 the value for using pointer optimized is smaller on all 3 platforms, specifically for windows on all 3 platforms.

6. CONCLUSION

Performing this project helped me develop a better understanding for the assembly language. It was helpful to me. Considering that my optimized code generated significant results in calculating the dot product using pointers optimized or indexes optimized, this means that I have a better understanding in the assembly language for 3 different platforms now. In general, it took some time to understand the procedure of what was happening before I could optimize the code but I was successful in doing it.